

Advanced Machine Learning

Final Project Presentation — MobileViTs for Sign Language Recognition



SAPIENZA
UNIVERSITÀ DI ROMA

Dario Loi (1940849)
Alessandro Monteleone (1883922)

M.Sc. in Computer Science,
Sapienza, University of Rome.

A. Y. 2024 – 2025


Task and Motivation

Task: Continuous Sign Language Recognition (CSLR). A **sequence-to-sequence** problem, where the input is a video and the output is a phrase in natural language.



Dataset:
RWTH-PHOENIX-Weather 2014T.
Around ≈ 4 thousand videos, 25 frames per second, 210×260 pixels.

MobileViTs

MobileViTs[2] are a lightweight version of Vision Transformers (ViTs) that are designed to run on mobile devices.



(a) Standard visual transformer (ViT)



(b) MobileViT. Here, $\text{Conv-}n \times n$ in the MobileViT block represents a standard $n \times n$ convolution and MV2 refers to MobileNetV2 block. Blocks that perform down-sampling are marked with $\downarrow 2$.

Figure: MobileViT architecture.

Our novel MobileViT architecture

We adapt the MobileViT architecture to CSLR by performing **spatio-temporal** aggregation across **multiple frames**. This comes at no additional cost in terms of parameters, while allowing full exploitation of the temporal dimension.




Figure: When the temporal dimension is processed in patches, the volume is called a **tubelet**.

First approach: Encoder-only MobileViT

To ensure a **minimal** amount of parameters, we first tried to train an encoder-only MobileViT. We use **CTC loss** to perform per-frame prediction of tokens from the encoder output, fusing repeated tokens.



Figure: Output of a CTC loss trained model. repeated tokens separated by a blank token are fused.

Second approach: Encoder-decoder MobileViT

To compare ourselves with a more classical approach, we also trained an **encoder-decoder** MobileViT. We attach a decoder head to the encoder output, and train the model with classic CE loss and teacher forcing.

This **massively** increases the number of parameters w.r.t. the encoder-only model, but it **simplifies** the training process.

Encoder-only MobileViT: Results

The encoder-only model gets trapped in a local minimum (predicting 0-length sequences) and does not recover. Repeated experiments with different configurations did **not** yield better results.



Figure: Validation loss for the encoder-only MobileViT.



Figure: Validation Word Error Rate for the encoder-only MobileViT, stuck at 1.0.

Encoder Decoder MobileViT: Results

The encoder-decoder approach shows a **smoother** convergence and a lower WER w.r.t. the encoder-only model.



Figure: Validation loss for the encoder-decoder MobileViT.



Figure: Validation Word Error Rate for the encoder-decoder MobileViT.

Quantitative results

Model	Parameters (M)	WER
Encoder-only XXS MobileViT (ours)	6.09	1.00
Encoder-decoder XXS MobileViT (ours)	7.80	0.79
SlowFastSign[1]	52.90	0.18
LCSA[3]	16.17	0.21




Table: Parameters (in millions) and validation Word-Error-Rate (WER) for the models.

Conclusions

CSLR is a challenging task that requires **massive** computational resources. Even while employing models that are **specifically** designed for mobile devices, we still struggle to achieve sufficient capacity to match state-of-the-art models.

For both the encoder-only and encoder-decoder models, limited experiments with increasing parameter size showed **improvements** in terms of WER. Showing that scaling the model up might be beneficial.

Thank you for your attention!
Questions?

-  J. Ahn, Y. Jang, and J. S. Chung.
SlowFast Network for Continuous Sign Language Recognition,
Sept. 2023.
[arXiv:2309.12304 \[cs\]](#).
-  S. Mehta and M. Rastegari.
MobileViT: Light-weight, General-purpose, and Mobile-friendly
Vision Transformer, Mar. 2022.
[arXiv:2110.02178 \[cs\]](#).
-  R. Zuo and B. Mak.
Local Context-aware Self-attention for Continuous Sign
Language Recognition.
In Interspeech 2022, pages 4810–4814. ISCA, Sept. 2022.